

Mushroom Classification

January 7, 2025

1 Introduction

1.1 Let's explore Mushrooms and try to classify them as poisonous or edible.

First we will perform EDA and try to get to grips with the data. Then we will train multiple classification models and compare their performance. Here are the models we will utilize: - Random Forest - Naive Bayes - Logistic Regression

```
[58]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, confusion_matrix, f1_score
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.metrics import roc_curve, roc_auc_score

[59]: # Loading the Mushroom dataset from the UCI ML repository
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/mushroom/
      ↪agaricus-lepiota.data"
column_names = ['class', 'cap-shape', 'cap-surface', 'cap-color', 'bruises',
                'odor', 'gill-attachment', 'gill-spacing', 'gill-size',
                'gill-color', 'stalk-shape', 'stalk-root',
                'stalk-surface-above-ring', 'stalk-surface-below-ring',
                'stalk-color-above-ring', 'stalk-color-below-ring', 'veil-type',
                'veil-color', 'ring-number', 'ring-type', 'spore-print-color',
                'population', 'habitat']

data = pd.read_csv(url, names=column_names)
data.head()

[59]:  class cap-shape cap-surface cap-color bruises odor gill-attachment \
0      p         x           s         n      t      p              f
1      e         x           s         y      t      a              f
```

2	e	b	s	w	t	l	f
3	p	x	y	w	t	p	f
4	e	x	s	g	f	n	f

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	\
0	c	n	k	...		s
1	c	b	k	...		s
2	c	b	n	...		s
3	c	n	n	...		s
4	w	b	k	...		s

	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	\
0		w		p	w
1		w		p	w
2		w		p	w
3		w		p	w
4		w		p	w

	ring-number	ring-type	spore-print-color	population	habitat
0	o	p		k	s u
1	o	p		n	n g
2	o	p		n	n m
3	o	p		k	s u
4	o	e		n	a g

[5 rows x 23 columns]

2 Lets do some EDA!

Here we will check for missing values, create a count plot to reveal if we have a balanced dataset, explore the feature distribution specifically: odor, cap-shape and gill-size.

The dataset is entirely categorical, therefore we need to preprocess the data and convert categorical data into numerical values that machine learning models can process.

The visualizations (count plots) for features like cap-shape, odor, and gill-size reveal how these features vary across the target classes. These patterns provide a better understanding of the dataset's structure and help confirm which features are most influential.

```
[60]: # Basic information about the dataset
data.info()

# Check for missing values
print("Missing values per column:")
print(data.isnull().sum())

# Countplot for the class distribution
custom_palette = ['#333652', '#FAD02C']
```

```

sns.countplot(x='class', data=data, hue='class', palette=custom_palette,
              legend=False)
plt.title('Distribution of Classes')
plt.xlabel('Class (Edible vs. Poisonous)')
plt.ylabel('Count')
plt.show()

# Explore the feature distributions
data.describe(include='all')

features_to_plot = ['odor', 'cap-shape', 'gill-size']
for feature in features_to_plot:
    sns.countplot(x=feature, data=data, hue='class', palette=custom_palette,
                  legend=False)
    plt.title(f'{feature} Distribution by Class')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.legend(title='Class', labels=['Edible', 'Poisonous'])
    plt.show()

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8124 entries, 0 to 8123
```


```
Data columns (total 23 columns):
```

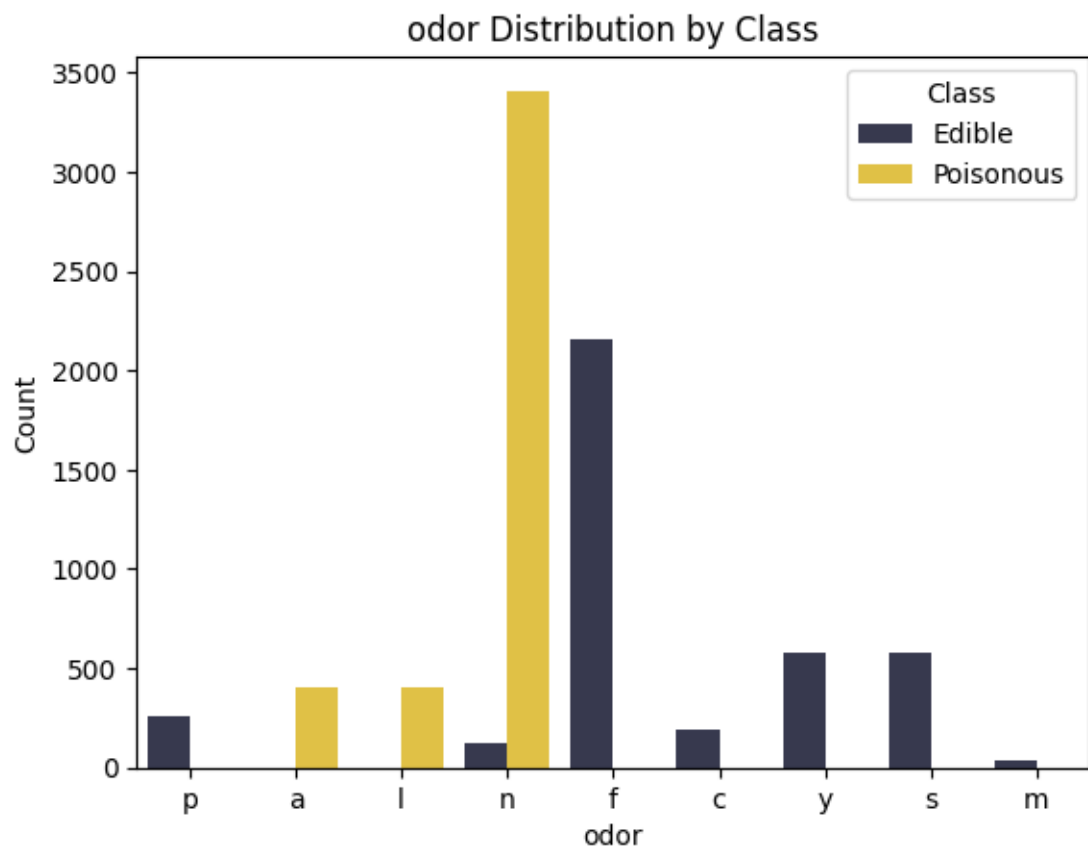
#	Column	Non-Null Count	Dtype
0	class	8124 non-null	object
1	cap-shape	8124 non-null	object
2	cap-surface	8124 non-null	object
3	cap-color	8124 non-null	object
4	bruises	8124 non-null	object
5	odor	8124 non-null	object
6	gill-attachment	8124 non-null	object
7	gill-spacing	8124 non-null	object
8	gill-size	8124 non-null	object
9	gill-color	8124 non-null	object
10	stalk-shape	8124 non-null	object
11	stalk-root	8124 non-null	object
12	stalk-surface-above-ring	8124 non-null	object
13	stalk-surface-below-ring	8124 non-null	object
14	stalk-color-above-ring	8124 non-null	object
15	stalk-color-below-ring	8124 non-null	object
16	veil-type	8124 non-null	object
17	veil-color	8124 non-null	object
18	ring-number	8124 non-null	object
19	ring-type	8124 non-null	object
20	spore-print-color	8124 non-null	object
21	population	8124 non-null	object

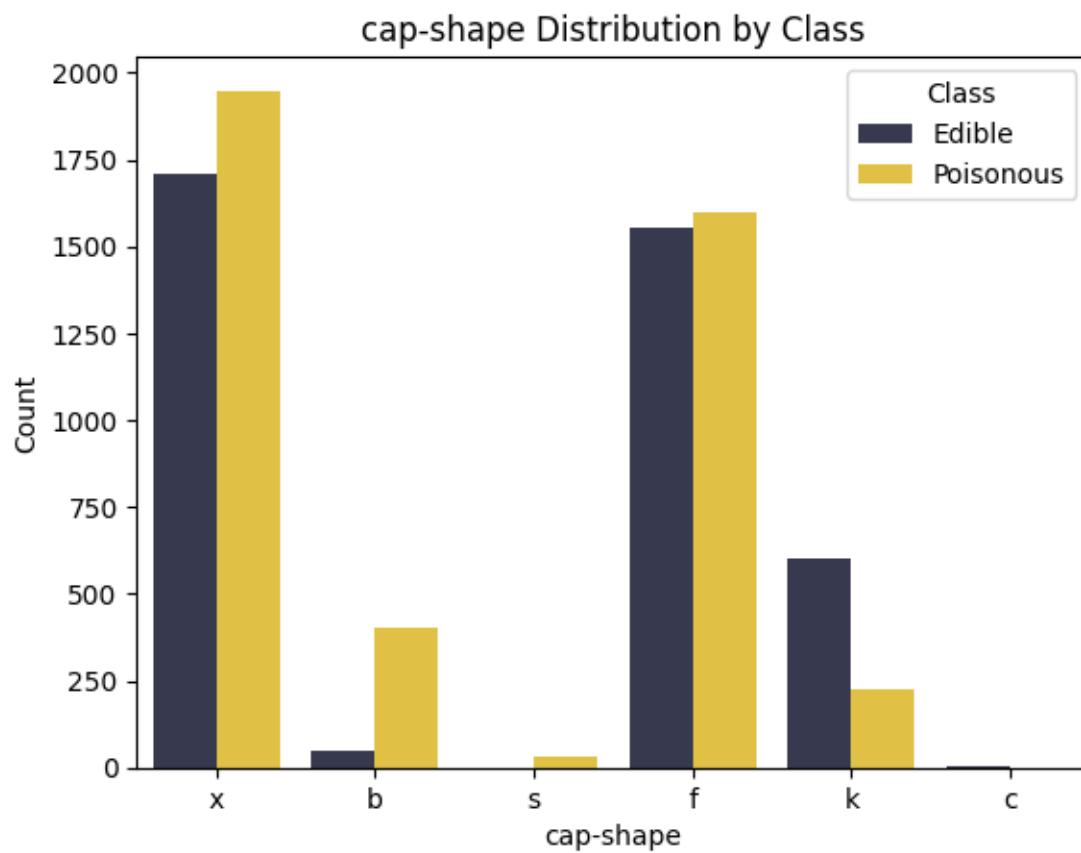
```

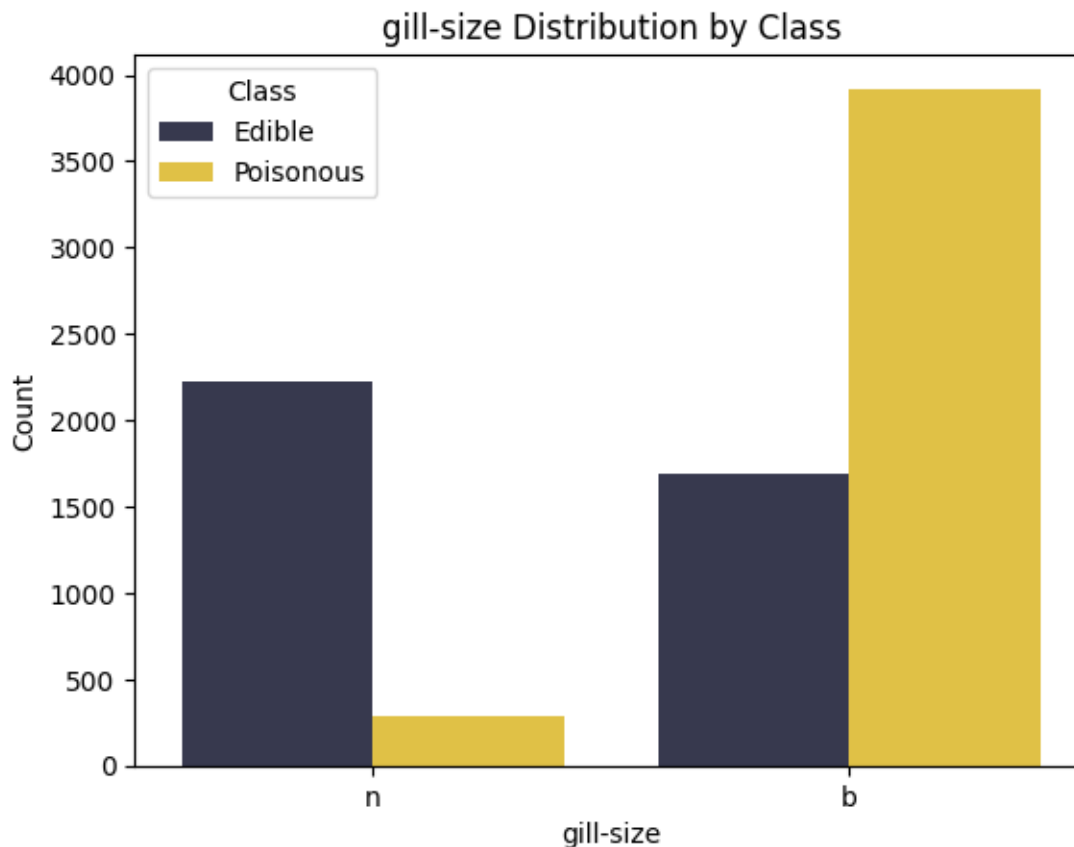
    22 habitat                                8124 non-null    object
dtypes: object(23)
memory usage: 1.4+ MB
Missing values per column:
class                                0
cap-shape                            0
cap-surface                          0
cap-color                            0
bruises                              0
odor                                  0
gill-attachment                      0
gill-spacing                         0
gill-size                            0
gill-color                           0
stalk-shape                          0
stalk-root                           0
stalk-surface-above-ring             0
stalk-surface-below-ring            0
stalk-color-above-ring              0
stalk-color-below-ring              0
veil-type                            0
veil-color                           0
ring-number                          0
ring-type                            0
spore-print-color                    0
population                           0
habitat                              0
dtype: int64

```









```
[61]: data.replace('?', np.nan, inplace=True)
data.dropna(inplace=True)
encoder = LabelEncoder()
for column in data.columns:
    data[column] = encoder.fit_transform(data[column])

print(data.info())
print(data.head())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 5644 entries, 0 to 8114
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	class	5644 non-null	int64
1	cap-shape	5644 non-null	int64
2	cap-surface	5644 non-null	int64
3	cap-color	5644 non-null	int64
4	bruises	5644 non-null	int64
5	odor	5644 non-null	int64

6	gill-attachment	5644	non-null	int64
7	gill-spacing	5644	non-null	int64
8	gill-size	5644	non-null	int64
9	gill-color	5644	non-null	int64
10	stalk-shape	5644	non-null	int64
11	stalk-root	5644	non-null	int64
12	stalk-surface-above-ring	5644	non-null	int64
13	stalk-surface-below-ring	5644	non-null	int64
14	stalk-color-above-ring	5644	non-null	int64
15	stalk-color-below-ring	5644	non-null	int64
16	veil-type	5644	non-null	int64
17	veil-color	5644	non-null	int64
18	ring-number	5644	non-null	int64
19	ring-type	5644	non-null	int64
20	spore-print-color	5644	non-null	int64
21	population	5644	non-null	int64
22	habitat	5644	non-null	int64

dtypes: int64(23)

memory usage: 1.0 MB

None

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	\
0	1	5	2	4	1	6	1	
1	0	5	2	7	1	0	1	
2	0	0	2	6	1	3	1	
3	1	5	3	6	1	6	1	
4	0	5	2	3	0	5	1	

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	\
0	0	1	2	...	2	
1	0	0	2	...	2	
2	0	0	3	...	2	
3	0	1	3	...	2	
4	1	0	2	...	2	

	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	\
0	5	5	0	0	
1	5	5	0	0	
2	5	5	0	0	
3	5	5	0	0	
4	5	5	0	0	

	ring-number	ring-type	spore-print-color	population	habitat
0	1	3	1	3	5
1	1	3	2	2	1
2	1	3	2	2	3
3	1	3	1	3	5
4	1	0	2	0	1

[5 rows x 23 columns]

3 Split and Train

Here we will split and train the data using 3 different ML models: - Random Forest - Naive Bayes - Logistic Regression

```
[62]: X = data.drop('class', axis=1)
      y = data['class']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                         random_state=42)
```

```
[63]: models = {
      "Random Forest": RandomForestClassifier(),
      "Naive Bayes": GaussianNB(),
      "Logistic Regression": LogisticRegression(max_iter=1000)
      }

      for name, model in models.items():
          print(f"Training {name}...")
          model.fit(X_train, y_train)
          y_pred = model.predict(X_test)

          print(f"{name} Classification Report:")
          print(classification_report(y_test, y_pred))

          print(f"Confusion Matrix for {name}:")
          print(confusion_matrix(y_test, y_pred))
          print("-" * 50)

      # Feature Importance for Random Forest
      rf_model = models["Random Forest"]
      feature_importances = pd.DataFrame({
          'Feature': X.columns,
          'Importance': rf_model.feature_importances_
      }).sort_values(by='Importance', ascending=False)

      custom_palette2 = ["#333652", "#FAD02C", "#4CAF50", "#00BFFF", "#4A4E69",
                          "#FFB400", "#3CB371", "#5DADE2", "#845EC2", "#FF6F91",
                          "#FFC75F", "#008F7A", "#C34A36", "#FF9671", "#2C73D2",
                          "#B39CD0", "#6A0572", "#D4AC0D", "#1E8449", "#3498DB",
                          "#E74C3C", "#F5B041"]

      plt.figure(figsize=(10, 6))
      sns.barplot(x='Importance', y='Feature', data=feature_importances,
                  hue = 'Feature', palette=custom_palette2, legend = False)
      plt.title('Feature Importance (Random Forest)')
```

```

plt.xlabel('Importance Score')
plt.ylabel('Feature')
plt.show()

# Predict probabilities for Logistic Regression
lr_model = models["Logistic Regression"]
y_pred_proba = lr_model.predict_proba(X_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = roc_auc_score(y_test, y_pred_proba)

plt.figure(figsize=(8, 6))
plt.plot(fpr,tpr,color='#FAD02C', marker='.', label=f'ROC AUC = {roc_auc:.2f}')
plt.title('ROC Curve (Logistic Regression)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid()
plt.show()

```

Training Random Forest...

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1058
1	1.00	1.00	1.00	636
accuracy			1.00	1694
macro avg	1.00	1.00	1.00	1694
weighted avg	1.00	1.00	1.00	1694

Confusion Matrix for Random Forest:

```

[[1058   0]
 [   0 636]]

```

Training Naive Bayes...

Naive Bayes Classification Report:

	precision	recall	f1-score	support
0	0.68	0.99	0.81	1058
1	0.94	0.23	0.37	636
accuracy			0.70	1694
macro avg	0.81	0.61	0.59	1694
weighted avg	0.78	0.70	0.64	1694

Confusion Matrix for Naive Bayes:

```
[[1049    9]
 [ 491 145]]
```


Training Logistic Regression...

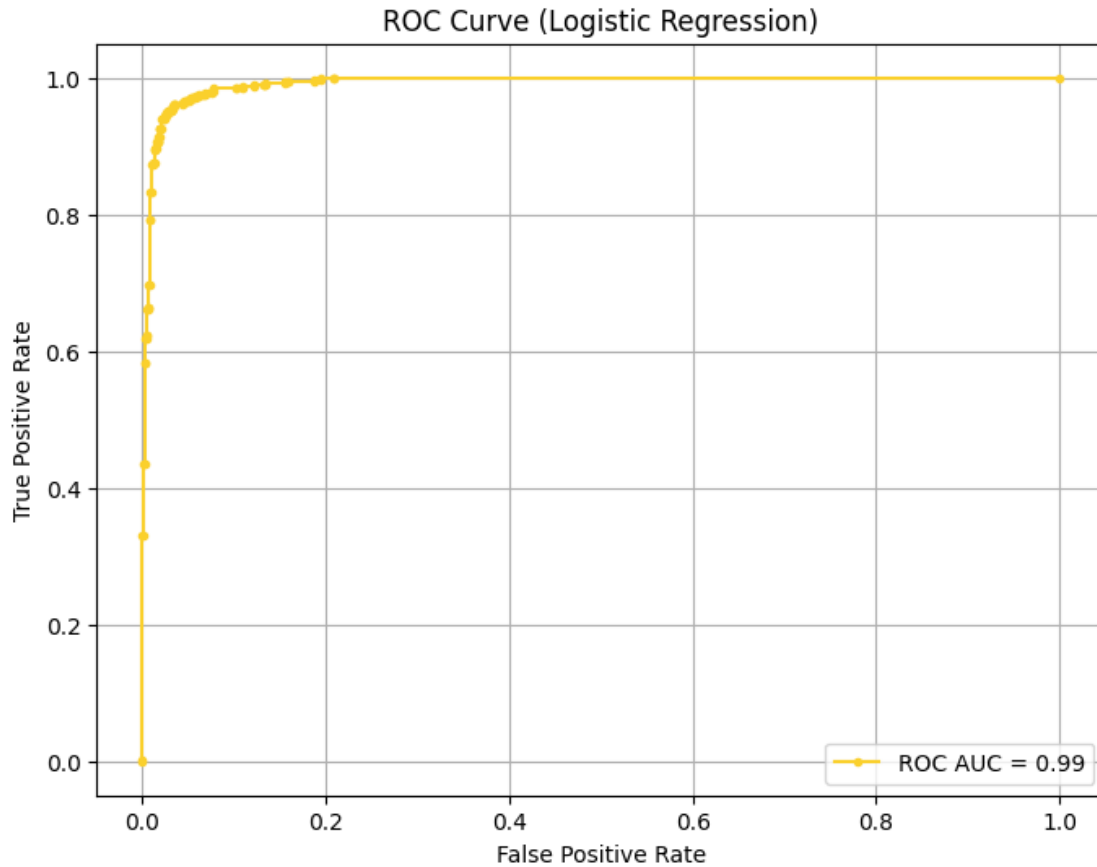
Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.96	0.98	0.97	1058
1	0.97	0.93	0.95	636
accuracy			0.96	1694
macro avg	0.96	0.96	0.96	1694
weighted avg	0.96	0.96	0.96	1694

Confusion Matrix for Logistic Regression:

```
[[1037   21]
 [  42 594]]
```






4 Time to Visualize!


In this section, we visualized the confusion matrices for each machine learning model to understand their performance in classifying mushrooms as edible or poisonous. Additionally, we compared the models using bar plots to evaluate their accuracy, precision, recall, and F1-score, highlighting the strengths and weaknesses of each algorithm.


```
[64]: def plot_confusion_matrix(cm, model_name):
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d',
                cmap=sns.light_palette("#333652", as_cmap=True), cbar=False)
    plt.title(f"Confusion Matrix: {model_name}")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

for name, model in models.items():
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
```

```
plot_confusion_matrix(cm, name)
```







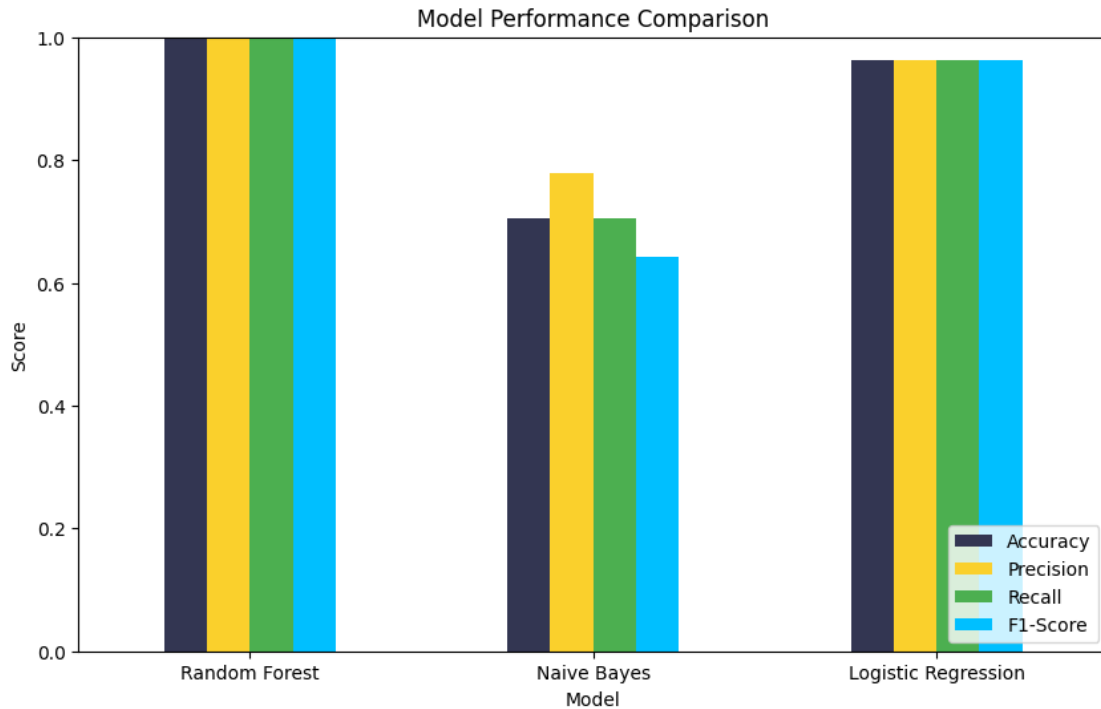
```
[65]: metrics = {  
    "Model": [],  
    "Accuracy": [],  
    "Precision": [],  
    "Recall": [],  
    "F1-Score": []  
}  
  
for name, model in models.items():  
    y_pred = model.predict(X_test)  
    report = classification_report(y_test, y_pred, output_dict=True)  
  
    metrics["Model"].append(name)  
    metrics["Accuracy"].append(report["accuracy"])  
    metrics["Precision"].append(report["weighted avg"]["precision"])  
    metrics["Recall"].append(report["weighted avg"]["recall"])  
    metrics["F1-Score"].append(report["weighted avg"]["f1-score"])  
  
metrics_df = pd.DataFrame(metrics)
```



```

custom_palette3 = ['#333652', '#FAD02C', '#4CAF50', '#00BFFF']
metrics_df.set_index("Model").plot(kind="bar", figsize=(10, 6),
                                   color=custom_palette3)
plt.title("Model Performance Comparison")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.ylim(0, 1)
plt.legend(loc="lower right")
plt.show()

```



5 Conclusion

The Random Forest Classifier outperformed the other models with a perfect score across all evaluation metrics: accuracy, precision, recall, and F1-score, each achieving 100%. This demonstrates its exceptional ability to handle categorical data effectively and capitalize on its ensemble approach, which reduces variance and prevents overfitting. The results confirm that Random Forest is well-suited for this dataset, likely due to its capacity to manage complex interactions between features and its robustness in classification tasks. In comparison, while Logistic Regression and Naive Bayes also performed well, they did not achieve the same level of precision and recall, highlighting Random Forest's advantage in leveraging feature importance to enhance classification accuracy.